

Average simulated dynamics under the optimal policy — complete method comparison at $p_0=0.1$

- Standard TS
- Scaled updates, $\eta=2$
- Temperature, $T=2$
- Observation forgetting, $\rho=0.95$
- Quality-maintenance mix
- Scaled updates, $\eta=0.5$
- Temperature, $T=0.5$
- Epsilon-greedy, $\epsilon=0.1$
- Calendar forgetting, $\rho=0.95$

