

Average simulated dynamics under the optimal policy — focused learning rules at  $p_0=0.3$

Standard TS   Scaled updates,  $\eta=2$    Epsilon-greedy,  $\epsilon=0.1$    Observation forgetting,  $\rho=0.95$    --- Quality-maintenance mix  
Scaled updates,  $\eta=0.5$

