

Average simulated dynamics under the optimal policy — focused learning rules at $p_0=0.5$

Standard TS Scaled updates, $\eta=2$ Epsilon-greedy, $\epsilon=0.1$ Observation forgetting, $\rho=0.95$ --- Quality-maintenance mix
Scaled updates, $\eta=0.5$

